



University of Pisa

Department of computer science

Algorithm design

Fifth hands-on: Bloom Filters

Domenico Erriquez

1. PROBLEM

1. Consider the Bloom filters where a single random universal hash random function $h: U \rightarrow [m]$ is employed for a set $S \subseteq U$ of keys, where U is the universe of keys.

Consider its binary array B of m bits. Suppose that $m \geq c|S|$, for some constant $c > 1$, and that both c and $|S|$ are unknown to us.

Estimate the expected number of 1s in B under a uniform choice at random of $h \in H$. Is this related to $|S|$? Can we use it to estimate $|S|$?

2. Consider B and its **rank** function: show how to use extra $O(m)$ bits to store a space-efficient data structure that returns, for any given i , the following answer in constant time:

$$\text{rank}(i) = \#1 \text{ in } B[1..i]$$

[Hint] Easy to solve in extra $O(m \log m)$ bits. To get $O(m)$ bits, use prefix sums on B , and sample them. Use a lookup table for pieces of B between any two consecutive samples.

2. SOLUTION FIRST QUESTION

In the first point, we aim to estimate the expected number of 1s in a bloom filter B with m positions using a single universal hash random function. To accomplish this, we define m indicator variables, where m is equal to the number of positions in B . Each indicator variable X_i $i \in [0 \dots m - 1]$ is defined as follows:

$$X_i = \begin{cases} 1, & \text{if } B[i] = 1 \\ 0, & \text{if } B[i] = 0 \end{cases}$$

We define a variable X as the sum of all indicator variables X_i . By calculating the expected value of X , we are able to estimate the number of 1s in the bloom filter B .

$$E[X] = E\left[\sum_{i=0}^{m-1} X_i\right] = E\sum_{i=0}^{m-1} E[X_i] = \sum_{i=0}^{m-1} \Pr[X_i = 1]$$

The probability that $B[i] = 0$ after n keys have been inserted is:

$$\Pr[B[i] = 0] = \left(1 - \frac{1}{m}\right)^n$$

So, we have the probability that that $B[i] = 1$ after n keys have been inserted is:

$$\Pr[B[i] = 1] = 1 - \Pr[B[i] = 0] = 1 - \left(1 - \frac{1}{m}\right)^n$$

Defining $|S| = n$, after n insertion the expected value of X is:

$$E[X] = m * \left(1 - \left(1 - \frac{1}{m}\right)^n\right) \approx m(1 - e^{-\frac{n}{m}})$$

The expected number of 1s in B is $m(1 - e^{-\frac{n}{m}})$ so it is strongly related to $n = |S|$

Defining $\mu = E[X] = m(1 - e^{-\frac{n}{m}})$ we can use it to estimate the $n = |S|$

$$\mu = m \left(1 - e^{-\frac{n}{m}}\right)$$

$$\frac{\mu}{m} = \left(1 - e^{-\frac{n}{m}}\right)$$

$$\frac{\mu}{m} - 1 = -e^{-\frac{n}{m}}$$

$$1 - \frac{\mu}{m} = e^{-\frac{n}{m}}$$

$$\ln\left(1 - \frac{\mu}{m}\right) = -\frac{n}{m}$$

$$n = -m \ln\left(1 - \frac{\mu}{m}\right)$$

3. SOLUTION SECOND QUESTION

The baseline solution is to precompute and store the prefix sum of each element of the binary array B , i.e., for each element we will store how many 1s occurs before it. To compute $rank(i)$, for any given i , the procedure just needs to return the i^{th} value of the precomputed prefix sums. Space complexity is $O(m \log m)$ bits because we store m prefix sums and the maximum prefix sum value (i.e., the maximum number of 1s) is m (when all the bits are set to 1) and we need $O(\log m)$ bits to store such value.

With the baseline solution is possible to answer $rank(i)$ in constant time, but we have to reduce the space required and reach the goal of $O(m)$ bits.

To improve space usage, instead of storing all prefix sums, we can store just a portion of them: we can subdivide B into blocks of length L and store the prefix sum of the block's ending position. With this approach less space is needed because less prefix sums are precomputed, but $rank(i)$ query takes time proportional to the block size which is more than constant time.

For example, having a binary array B with $m = 16$ positions, and the length of the blocks $L = \text{Log}_2 m = 4$, the array P containing the prefix sums of the blocks is the following:

$$\begin{array}{c}
 \downarrow \quad \downarrow \quad \downarrow \quad \downarrow \\
 B = [0,0,1,0,0,1,0,1,0,1,0,0,1,0,1,1] \\
 P = [1,3,4,7]
 \end{array}$$

To improve time complexity, it would be ideal to also precompute the prefix sum of any position inside the block. For example, given a block $[0 \ 1 \ 0 \ 1]$, we would like to know in advance that, from left to right:

- at the first position, the prefix sum is 0.
- at the second, the prefix sum is 1.
- at the third, the prefix sum 1 as well.
- at the last position, the prefix sum is 2.

So to have constant time complexity we need to find a way to retrieve these information in constant time. A possible solution is to create a lookup table where we save for each row of the table the precomputed prefix sum of any possible block, and since each block is a combination of 0s and 1s it is an integer number and can be used to index to its precomputed prefix sum.

For example, having a binary array B with $m = 16$ positions, and the length of the blocks $L = \text{Log}_2 m = 4$, we have the following lookup table T :

| | | | | |
|------------|-----|-----|-----|-----|
| $(0000)_2$ | 0 | 0 | 0 | 0 |
| ... | ... | ... | ... | ... |
| $(0011)_2$ | 0 | 0 | 1 | 2 |
| ... | ... | ... | ... | ... |
| $(0100)_2$ | 0 | 1 | 1 | 1 |
| ... | ... | ... | ... | ... |
| $(1011)_2$ | 1 | 1 | 2 | 3 |
| ... | ... | ... | ... | ... |
| $(1111)_2$ | 1 | 2 | 3 | 4 |

The number of all possible blocks of length L is 2^L . So, the table will have 2^L rows and L columns. Finally, the solution to answer a query $\text{rank}(i)$ in constant time works like this:

$$\text{rank}(i) = \begin{cases} P \left\lceil \frac{i+1}{L} \right\rceil - 1, & \text{if } i+1 \% L == 0 \\ P \left\lceil \frac{i+1}{L} \right\rceil + T[B[L * C, L * C + L - 1]][i \% L], & \text{otherwise} \end{cases}$$

Where $C = \left\lceil \frac{i+1}{L} \right\rceil$

3.1. Space complexity

The space used is the space needed to store the block's prefix sums and the table. The number of blocks is $\frac{m}{L}$ and the bits needed to store a prefix sum is $O(\log m)$ so space needed to store blocks' prefix sum is $O(\frac{m}{L} * \log m)$ bits. The table has 2^L rows and L columns and an element of a cell can be stored in $O(\log L)$ bits, so table's space complexity is $O(2^L * L * \log L)$ bits.